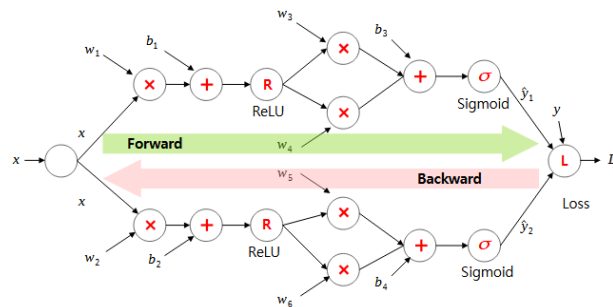


# GradientTape

## 4. TensorFlow & Keras

### Part 1: Ways to build neural networks in TensorFlow



This video was produced in Korean and translated into English,  
and the audio was generated by AI (TTS).

[www.youtube.com/@meanxai](http://www.youtube.com/@meanxai)

## TensorFlow / Keras

- [MXDL-4-01] {
  - 1. Multiple ways to build neural networks in TensorFlow and Keras API
  - 2. Type-1: Tensorflow's GradientTape() & gradient descent – [Example] Binary Classification
  - 3. Type-2: Tensorflow's GradientTape() & Optimizer – [Example] Multiclass Classification
  - 4. Type-3: Tensorflow's Optimizer – [Example] Nonlinear Regression
- [MXDL-4-02] {
  - 5. Type-4: Keras Sequential model – [Example] Binary Classification
  - 6. Type-5: Keras Functional API - [Example] Multiclass Classification
  - 7. Type-6: Tensorflow + Keras Functional API - [Example] Nonlinear Regression
  - 8. Type-7: Customizing Keras - [Example] Custom loss (Regularized loss function)


## Multiple ways to build neural networks in TensorFlow and Keras API

Create a neural network

```
wh = tf.Variable(np.random.normal(size=(ni, nh)))
bh = tf.Variable(np.zeros(shape=(1, nh)))
wo = tf.Variable(np.random.normal(size=(nh, no)))
bo = tf.Variable(np.zeros(shape=(1, no)))
parameters = [wh, bh, wo, bo]
```

```
def predict(x): # forward propagation
    p = parameters
    h = tf.nn.relu(tf.matmul(x, p[0]) + p[1])
    y_hat = tf.sigmoid(tf.matmul(h, p[2]) + p[3])
    return y_hat
```

$$\hat{y}(w_h, b_h, w_o, b_o) = \sigma(R(x \cdot w_h + b_h) \cdot w_o + b_o)$$



### 1. GradientTape & Gradient descent

```
with tf.GradientTape() as tape:
    loss = Loss(y, y_hat)

grads = tape.gradient(loss, parameters)

for i, p in enumerate(parameters):
    p.assign_sub(lr * grads[i])
```

### 2. GradientTape & Optimizer

```
opt = optimizers.Adam(learning_rate=0.01,
                       beta_1=0.9, beta_2=0.999)
```

```
with tf.GradientTape() as tape:
    loss = Loss(y, y_hat)

grads = tape.gradient(loss, parameters)

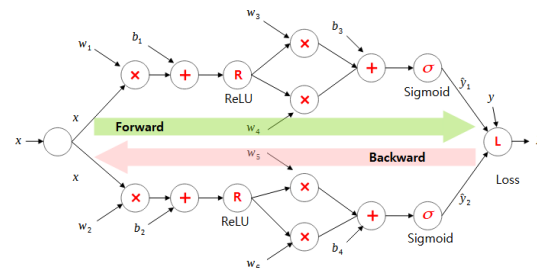
opt.apply_gradients(zip(grads, parameters))
```

### 3. Optimizer

```
opt = optimizers.Adam(learning_rate = 0.01)

loss = Loss(y, y_hat)
opt.minimize(loss, parameters)
```

### Automatic Differentiation



### Adam optimizer

$$m_t = \beta m_{t-1} + (1-\beta)g_t$$

$$G_t = \rho G_{t-1} + (1-\rho)g_t^2$$

$$\hat{m}_t = \frac{m_t}{1-\beta^n} \quad \hat{G}_t = \frac{G_t}{1-\rho^n}$$

$$w_t = w_{t-1} - \frac{\alpha}{\sqrt{\hat{G}_t + \epsilon}} \cdot \hat{m}_t$$

## ■ Multiple ways to Build Neural Networks in TensorFlow and Keras API

### 4. Sequential Model

```
model = Sequential()
model.add(Dense(ni, input_dim=n1, activation='relu'))
model.add(Dense(no, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam')

h = model.fit(x, y, epochs=200, batch_size=50)
```

### 5. Functional API Model

```
x_input = Input(batch_shape=(None, ni))
h = Dense(nh, activation='relu')(x_input)
y_output = Dense(no, activation='sigmoid')(h)
model = Model(x_input, y_output)
model.compile(loss='binary_crossentropy', optimizer='adam')


h = model.fit(x, y, epochs=200, batch_size=50)
```

### 6. Functional API Model + Tensorflow

```
x_input = Input(batch_shape=(None, ni))
h = Dense(nh, activation='relu')(x_input)
y_output = Dense(no, activation='sigmoid')(h)
model = Model(x_input, y_output)

for i in range(epochs):
    with tf.GradientTape() as tape:
        loss = Loss(y, y_hat)

    grads = tape.gradient(loss, model.trainable_variables)
    opt.apply_gradients(zip(grads, model.trainable_variables))
```



### 7. Customizing


- Customizing Layer
- Customizing Loss function
- Customizing Model fit()

■ **Type-1.** Tensorflow's GradientTape & Gradient descent – [Example] Binary Classification

```
# [MXDL-4-01] 1.tf_binary_class.py
# Binary classification
import numpy as np
import tensorflow as tf
from sklearn.datasets import make_blobs
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# Generate a data set
x, y = make_blobs(n_samples=300, n_features=2,
                  centers=[[0., 0.], [0.5, 0.1]],
                  cluster_std=0.2, center_box=(-1., 1.))


y = y.reshape(-1,1)
x_train, x_test, y_train, y_test = train_test_split(x, y)
```



```
# Visually see the data.
plt.figure(figsize=(6,4))
color = [['red', 'blue'][a] for a in y.reshape(-1,)]
plt.scatter(x[:,0], x[:,1],s=100,c=color,alpha=0.3)
plt.show()
```

```
# Create an ANN with a hidden layer
n_input = x.shape[1] # number of input neurons
n_output = 1         # number of output neurons
n_hidden = 8         # number of hidden neurons
lr = 0.05            # learning rate
```

```
# Initialize the parameters
wh = tf.Variable(np.random.normal(size=(n_input, n_hidden)))
bh = tf.Variable(np.zeros(shape=(1, n_hidden)))
wo = tf.Variable(np.random.normal(size=(n_hidden, n_output)))
bo = tf.Variable(np.zeros(shape=(1, n_output)))
parameters = [wh, bh, wo, bo]
```



```
# loss function
def binary_crossentropy(y, y_hat):
    return -tf.reduce_mean(y * tf.math.log(y_hat) + \
                           (1. - y) * tf.math.log(1. - y_hat))

def predict(x, proba=True):
    p = parameters
    o_hidden = tf.nn.relu(tf.matmul(x, p[0]) + p[1])
    o_output = tf.sigmoid(tf.matmul(o_hidden, p[2]) + p[3])

    if proba:
        return o_output # return sigmoid output as is
    else:
        return (o_output.numpy() > 0.5) * 1 # return class

def fit(x_trn, y_trn, x_val, y_val, epochs, batch_size):
    trn_loss = []
    val_loss = []
    for epoch in range(epochs):
        # Training with mini-batch
        for batch in range(int(x_trn.shape[0] / batch_size)):
            idx = np.random.choice(x_trn.shape[0], batch_size)
            x_bat = x_trn[idx]
            y_bat = y_trn[idx]

            # Automatic differentiation
            with tf.GradientTape() as tape:
                loss = binary_crossentropy(y_bat, predict(x_bat))

            # Find the gradients of loss w.r.t the parameters
            grads = tape.gradient(loss, parameters)

            # update parameters by the gradient descent
            for i, p in enumerate(parameters):
                p.assign_sub(lr * grads[i]) # p = p - lr * gradient
```

## ■ Type-1. Tensorflow's GradientTape & Gradient descent – [Example] Binary Classification

```
# loss history
loss = binary_crossentropy(y_trn, predict(x_trn))
trn_loss.append(loss.numpy())

loss = binary_crossentropy(y_val, predict(x_val))
val_loss.append(loss.numpy())

if epoch % 10 == 0:
    print("{}: train_loss={:.4f}, val_loss={:.4f}".\
          format(epoch, trn_loss[-1], val_loss[-1]))

return trn_loss, val_loss
```


```
# training
trn_loss, val_loss = fit(x_train, y_train, x_test, y_test,
                        epochs=200, batch_size=50)
```

```
# Visually see the loss history
plt.plot(trn_loss, c='blue', label='train loss')
plt.plot(val_loss, c='red', label='validation loss')
plt.legend()
plt.show()
```

```
# Check the accuracy of the test data
y_pred = predict(x_test, proba=False)
acc = (y_pred == y_test).mean()
print("\nAccuracy of the test data = {:.4f}".format(acc))
```

```
0: train_loss=0.6704, val_loss=0.6838
10: train_loss=0.4978, val_loss=0.5131
20: train_loss=0.4015, val_loss=0.4194
30: train_loss=0.3399, val_loss=0.3568
40: train_loss=0.3006, val_loss=0.3195
50: train_loss=0.2754, val_loss=0.2931
60: train_loss=0.2590, val_loss=0.2813
70: train_loss=0.2466, val_loss=0.2599
80: train_loss=0.2374, val_loss=0.2559
90: train_loss=0.2305, val_loss=0.2448
100: train_loss=0.2254, val_loss=0.2425
110: train_loss=0.2216, val_loss=0.2319
120: train_loss=0.2176, val_loss=0.2284
...
160: train_loss=0.2085, val_loss=0.2196
170: train_loss=0.2078, val_loss=0.2275
180: train_loss=0.2062, val_loss=0.2242
190: train_loss=0.2052, val_loss=0.2211
```

Accuracy of the test data = 0.92



## ■ Type-2. Tensorflow's GradientTape & Optimizer – [Example] Multiclass Classification

# [MXDL-4-01] 2.tf\_multi\_class.py: Multiclass classification

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import optimizers
from sklearn.datasets import make_blobs
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

# Generate a dataset

```
x, y = make_blobs(n_samples=400, n_features=2,
                  centers=[[0., 0.], [0.5, 0.1], [1., 0.]],
                  cluster_std=0.15, center_box=(-1., 1.))
n_class = np.unique(y).shape[0] # the number of classes
```

# one-hot encode class y, y = [0,1,2]

```
y_ohe = np.eye(n_class)[y]
x_train, x_test, y_train, y_test = train_test_split(x, y_ohe)
```

# Visually see the data.



```
plt.figure(figsize=(5,4))
color = [['red', 'blue', 'green']][a] \
        for a in y.reshape(-1,)
plt.scatter(x[:,0], x[:,1], s=100, c=color, alpha=0.3)
plt.show()
```

# Create an ANN with a hidden layer

```
n_input = x.shape[1] # number of input neurons
n_output = n_class # number of output neurons
n_hidden = 8 # number of hidden neurons
lr = 0.05 # learning rate
```

# Initialize the parameters

```
wh = tf.Variable(np.random.normal(size=(n_input, n_hidden)))
bh = tf.Variable(np.zeros(shape=(1, n_hidden)))
wo = tf.Variable(np.random.normal(size=(n_hidden, n_output)))
bo = tf.Variable(np.zeros(shape=(1, n_output)))
parameters = [wh, bh, wo, bo]
```



```
opt = optimizers.Adam(learning_rate=0.01, beta_1=0.9, beta_2=0.999)
```

Adam optimizer

# loss function

```
def crossentropy(y, y_hat):
    ce = -tf.reduce_sum(y * tf.math.log(y_hat),
                        axis=1)
    return tf.reduce_mean(ce)
```

def predict(x, proba=True):

```
p = parameters
o_hidden = tf.nn.relu(tf.matmul(x, p[0]) + p[1])
o_output = tf.nn.softmax(tf.matmul(o_hidden, p[2]) + p[3])
```

if proba:

```
    return o_output # return softmax output as is
```

else:

```
    return tf.math.argmax(o_output, axis=1) # return class
```

def fit(x\_trn, y\_trn, x\_val, y\_val, epochs, batch\_size):

```
    trn_loss, val_loss = [], []
    for epoch in range(epochs):
```

# Training with mini-batch

```
    for batch in range(int(x_trn.shape[0] / batch_size)):
        idx = np.random.choice(x_trn.shape[0], batch_size)
        x_batch = x_trn[idx]
        y_batch = y_trn[idx]
```

# Automatic differentiation

```
    with tf.GradientTape() as tape:
        loss = crossentropy(y_batch, predict(x_batch))
```

# Find the gradients and update the parameters

```
    grads = tape.gradient(loss, parameters)
    opt.apply_gradients(zip(grads, parameters))
```

$$m_t = \beta m_{t-1} + (1-\beta) g_t$$

$$G_t = \rho G_{t-1} + (1-\rho) g_t^2$$

$$\hat{m}_t = \frac{m_t}{1-\beta^t} \quad \hat{G}_t = \frac{G_t}{1-\rho^t}$$

$$w_t = w_{t-1} - \frac{\alpha}{\sqrt{\hat{G}_t + \epsilon}} \hat{m}_t$$

## ■ Type-2. Tensorflow's GradientTape & Optimizer – [Example] Multiclass Classification

```
# loss history
loss = crossentropy(y_trn, predict(x_trn))
trn_loss.append(loss.numpy())

loss = crossentropy(y_val, predict(x_val))
val_loss.append(loss.numpy())

if epoch % 10 == 0:
    print("{}: train_loss={:.4f}, val_loss={:.4f}".\
          format(epoch, trn_loss[-1], val_loss[-1]))

return trn_loss, val_loss


# training
train_loss, test_loss = fit(x_train, y_train, x_test, y_test,
                           epochs=200, batch_size=50)

# Visually see the loss history
plt.plot(train_loss, c='blue', label='train loss')
plt.plot(test_loss, c='red', label='test loss')
plt.legend()
plt.show()

# Check the accuracy of the test data
y_pred = predict(x_test, proba=False).numpy()
acc = (y_pred == np.argmax(y_test, axis=1)).mean()
print("Accuracy of test data = {:.4f}".format(acc))
```

```
0: train_loss=1.1298, val_loss=1.1364
10: train_loss=0.3957, val_loss=0.4014
20: train_loss=0.2158, val_loss=0.2312
30: train_loss=0.1754, val_loss=0.1883
40: train_loss=0.1559, val_loss=0.1764
50: train_loss=0.1465, val_loss=0.1847
60: train_loss=0.1371, val_loss=0.1708
70: train_loss=0.1333, val_loss=0.1656
80: train_loss=0.1306, val_loss=0.1596
90: train_loss=0.1282, val_loss=0.1748
100: train_loss=0.1285, val_loss=0.1864
110: train_loss=0.1243, val_loss=0.1652
120: train_loss=0.1223, val_loss=0.1607
...
160: train_loss=0.1198, val_loss=0.1816
170: train_loss=0.1194, val_loss=0.1623
180: train_loss=0.1204, val_loss=0.1793
190: train_loss=0.1233, val_loss=0.1569
```

Accuracy of the test data = 0.92





## ■ Type-3. Tensorflow's Optimizer – [Example] Nonlinear Regression

# [MXDL-4-01] 3.tf\_regression.py Nonlinear Regression

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import optimizers
from sklearn.datasets import make_blobs
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

# Generate a data set

```
x = np.random.random((1000, 1))
y = 2.0 * np.sin(2.0 * np.pi * x) + \
    np.random.normal(0.0, 0.8, (1000, 1))
```

# Generate training, test data set



```
x_train, x_test, y_train, y_test = train_test_split(x, y)
x_pred = np.linspace(0, 1, 200).reshape(-1, 1)
```

# Visually see the data.

```
plt.figure(figsize=(7,5))
plt.scatter(x_train, y_train, s=20, c='blue', alpha=0.3, label='train')
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.3, label='valid')
plt.legend()
plt.show()
```

# Create an ANN with a hidden layer

```
n_input = x.shape[1] # number of input neurons
n_output = 1         # number of output neurons
n_hidden = 8         # number of hidden neurons
lr = 0.05            # learning rate
```



# Initialize the parameters

```
wh = tf.Variable(np.random.normal(size=(n_input, n_hidden)))
bh = tf.Variable(np.zeros(shape=(1, n_hidden)))
wo = tf.Variable(np.random.normal(size=(n_hidden, n_output)))
bo = tf.Variable(np.zeros(shape=(1, n_output)))
parameters = [wh, bh, wo, bo]
opt = optimizers.Adam(learning_rate = 0.01)
```

# loss function: mean squared error

```
def mse(y, y_hat):
    return tf.reduce_mean(tf.math.square(y - y_hat))
```

def predict(x, proba=True):

```
    p = parameters
    o_hidden = tf.math.tanh(tf.matmul(x, p[0]) + p[1])
    o_output = tf.matmul(o_hidden, p[2]) + p[3]
    return o_output
```

def fit(x\_trn, y\_trn, x\_val, y\_val, epochs, batch\_size):

```
    trn_loss = []
    val_loss = []
    for epoch in range(epochs):
        # Training with mini-batch
        for batch in range(int(x_trn.shape[0] / batch_size)):
            idx = np.random.choice(x_trn.shape[0], batch_size)
            x_bat = x_trn[idx]
            y_bat = y_trn[idx]
```

# Automatic differentiation and update parameters

```
    loss = lambda: mse(y_bat, predict(x_bat))
    opt.minimize(loss, parameters)
```

## ■ Type-3. Tensorflow's Optimizer – [Example] Nonlinear Regression

```
# loss history
loss = mse(y_trn, predict(x_trn))
trn_loss.append(loss.numpy())

loss = mse(y_val, predict(x_val))
val_loss.append(loss.numpy())

if epoch % 10 == 0:
    print("{}: train_loss={:.4f}, val_loss={:.4f}".\
          format(epoch, trn_loss[-1], val_loss[-1]))

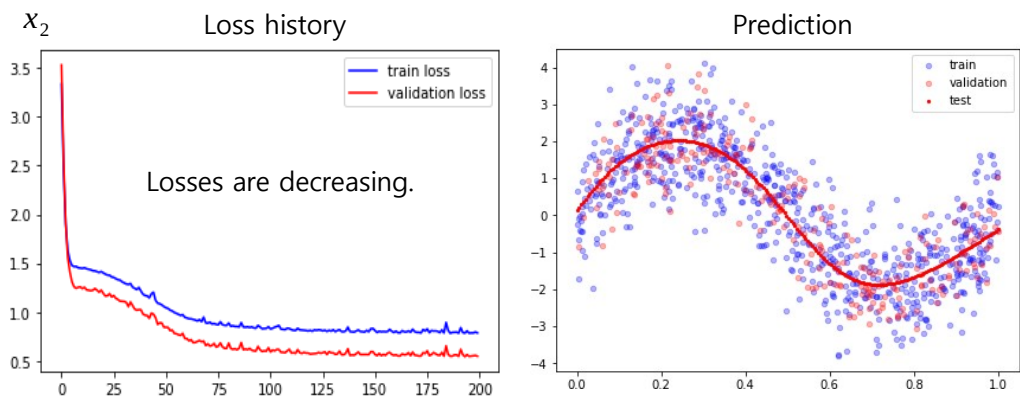
return trn_loss, val_loss


# training
trn_loss, val_loss = fit(x_train, y_train, x_test, y_test,
                        epochs=200, batch_size=50)

# Visually see the loss history.
plt.plot(trn_loss, c='blue', label='train loss')
plt.plot(val_loss, c='red', label='validation loss')
plt.legend()
plt.show()

# Visually check the prediction result.
y_pred = predict(x_pred)
plt.figure(figsize=(7,5))
plt.scatter(x_train, y_train, s=20, c='blue', alpha=0.3,
            label='train')
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.3,
            label='validation')
plt.scatter(x_pred, y_pred, s=5, c='red', label='test')
plt.legend()
plt.show()
```

```
0: train_loss=3.3347, val_loss=3.5312
10: train_loss=1.4471, val_loss=1.2408
20: train_loss=1.4016, val_loss=1.1910
30: train_loss=1.2989, val_loss=1.1085
40: train_loss=1.1695, val_loss=0.9670
50: train_loss=1.0439, val_loss=0.8442
60: train_loss=0.9381, val_loss=0.7275
70: train_loss=0.8866, val_loss=0.6739
80: train_loss=0.8866, val_loss=0.6811
90: train_loss=0.8418, val_loss=0.6271
100: train_loss=0.8269, val_loss=0.5914
110: train_loss=0.8194, val_loss=0.5796
120: train_loss=0.8037, val_loss=0.5726
130: train_loss=0.8178, val_loss=0.5819
140: train_loss=0.7953, val_loss=0.5569
150: train_loss=0.8094, val_loss=0.5655
160: train_loss=0.8048, val_loss=0.5616
170: train_loss=0.7917, val_loss=0.5553
180: train_loss=0.8019, val_loss=0.5596
190: train_loss=0.7834, val_loss=0.5459
```





## 4. TensorFlow & Keras

### Part 2: Ways to build neural networks in Keras

This video was produced in Korean and translated into English,  
and the audio was generated by AI (TTS).

[www.youtube.com/@meanxai](http://www.youtube.com/@meanxai)

## ■ Multiple ways to Build Neural Networks in Keras

### 1. Sequential API

```
model = Sequential()
model.add(Dense(n1, input_dim=n1, activation='relu'))
model.add(Dense(n2, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam')

h = model.fit(x, y, epochs=200, batch_size=50)
```

### 2. Functional API

```
x_input = Input(batch_shape=(None, n1))
h = Dense(n2, activation='relu')(x_input)
y_output = Dense(n3, activation='sigmoid')(h)
model = Model(x_input, y_output)
model.compile(loss='binary_crossentropy', optimizer='adam')


h = model.fit(x, y, epochs=200, batch_size=50)
```

### 3. Functional API + Tensorflow

```
x_input = Input(batch_shape=(None, n1))
h = Dense(n2, activation='relu')(x_input)
y_output = Dense(n3, activation='sigmoid')(h)
model = Model(x_input, y_output)

for i in range(epochs):
    with tf.GradientTape() as tape:
        loss = Loss(y, y_hat)

    grads = tape.gradient(loss, model.trainable_variables)
    opt.apply_gradients(zip(grads, model.trainable_variables))
```



### 4. Customizing

- Custom Loss function
- Custom Layer
- Custom Model fit()

## 1. Keras Sequential API – [Example] Binary Classification

# [MXDL-4-02] 4.keras\_binary\_class.py - Binary classification

```
import numpy as np
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras import optimizers
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
```

# Generate a data set

```
x, y = make_blobs(n_samples=300, n_features=2,
                  centers=[[0., 0.], [0.5, 0.1]],
                  cluster_std=0.2, center_box=(-1., 1.))
y = y.reshape(-1,1)
x_train, x_test, y_train, y_test = train_test_split(x, y)
```

# Visually see the data.



```
plt.figure(figsize=(6,4))
color = [['red', 'blue'][a] for a in y.reshape(-1,)]
plt.scatter(x[:,0], x[:,1], s=100, c=color, alpha=0.3)
plt.show()
```

# Create an ANN with a hidden layer

```
n_input = x.shape[1] # number of input neurons
n_output = 1         # number of output neurons
n_hidden = 8         # number of hidden neurons
adam = optimizers.Adam(learning_rate=0.01)
```

# Create an ANN model

```
model = Sequential()
model.add(Dense(n_hidden, input_dim=n_input, activation='relu'))
model.add(Dense(n_output, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer=adam)
```



# training

```
f = model.fit(x_train, y_train,
              validation_data=(x_test, y_test),
              epochs=200, batch_size=50)
```

# Visually see the loss history

```
plt.plot(f.history['loss'], c='blue', label='train loss')
plt.plot(f.history['val_loss'], c='red', label='validation loss')
plt.legend()
plt.show()
```

# Check the accuracy of the test data

```
y_pred = (model.predict(x_test) > 0.5) * 1
acc = (y_pred == y_test).mean()
print("\nAccuracy of the test data = {:.2f}".format(acc))
```

Epoch 198/200

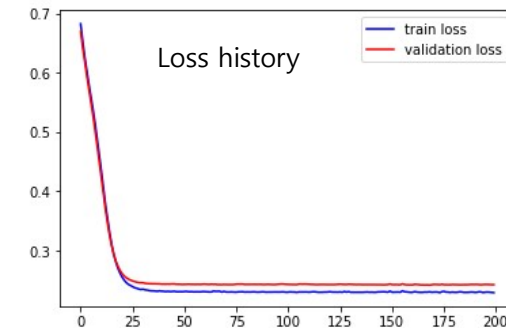
5/5 [=====] - 0s 10ms/step - loss: 0.2377 - val\_loss: 0.1801

Epoch 199/200

5/5 [=====] - 0s 6ms/step - loss: 0.2377 - val\_loss: 0.1799

Epoch 200/200

5/5 [=====] - 0s 7ms/step - loss: 0.2376 - val\_loss: 0.1806



Accuracy of the test data = 0.92

## ■ 2. Keras functional API – [Example] Multiclass Classification

# [MXDL-4-02] 5.keras\_multi\_class.py - Multiclass classification

```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.datasets import make_blobs
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras import optimizers
import matplotlib.pyplot as plt
```

# Generate a dataset for multiclass classification

```
x, y = make_blobs(n_samples=400, n_features=2,
                  centers=[[0., 0.], [0.5, 0.1], [1., 0.]],
                  cluster_std=0.15, center_box=(-1., 1.))
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y)
n_class = np.unique(y).shape[0] # the number of classes
```

# Create an ANN with a hidden layer



```
n_input = x.shape[1] # number of input neurons
n_output = n_class # number of output neurons
n_hidden = 8 # number of hidden neurons
adam = optimizers.Adam(learning_rate=0.01)
```

# Create an ANN model

```
x_input = Input(batch_shape=(None, n_input))
h = Dense(n_hidden, activation='relu')(x_input)
y_output = Dense(n_output, activation='softmax')(h)
model = Model(x_input, y_output)
model.compile(loss='sparse_categorical_crossentropy',
              optimizer=adam)
```

# training

```
f = model.fit(x_train, y_train,
              validation_data=(x_test, y_test),
              epochs=200, batch_size=50)
```



# Visually see the data.

```
plt.figure(figsize=(5,4))
color = [['red', 'blue', 'green'][a] for a in y.reshape(-1,)]
plt.scatter(x[:, 0], x[:, 1], s=70, c=color, alpha=0.3)
plt.show()
```

# Visually see the loss history

```
plt.plot(f.history['loss'], c='blue', label='train loss')
plt.plot(f.history['val_loss'], c='red', label='validation loss')
plt.legend()
plt.show()
```

# Check the accuracy of the test data

```
y_pred = model.predict(x_test)
acc = (np.argmax(y_pred, axis=1) == y_test).mean()
print("Accuracy of test data = {:.2f}".format(acc))
```

Epoch 198/200


6/6 [=====] - 0s 4ms/step - loss: 0.1207 - val\_loss: 0.1395

Epoch 199/200

6/6 [=====] - 0s 4ms/step - loss: 0.1210 - val\_loss: 0.1411

Epoch 200/200

6/6 [=====] - 0s 4ms/step - loss: 0.1209 - val\_loss: 0.1385



Accuracy of the test data = 0.96

### 3. Tensorflow + Keras functional API – [Example] Nonlinear Regression

# [MXDL-4-02] 6.tf\_keras\_regression.py - Nonlinear Regression

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import make_blobs
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras import optimizers
import matplotlib.pyplot as plt
```

# Generate a data set

```
x = np.random.random((1000, 1))
y = 2.0 * np.sin(2.0 * np.pi * x) + \
    np.random.normal(0.0, 0.8, (1000, 1))
```

# Generate training, test data set

```
x_train, x_test, y_train, y_test = train_test_split(x, y)
x_pred = np.linspace(0, 1, 200).reshape(-1, 1)
```



# loss function: mean squared error

```
def mse(y, y_hat):
    return tf.reduce_mean(tf.math.square(y-y_hat))
```

# Create an ANN model

```
n_input = x.shape[1] # number of input neurons
n_output = 1         # number of output neurons
n_hidden = 8         # number of hidden neurons
opt = optimizers.Adam(learning_rate=0.01)
```

```
x_input = Input(batch_shape=(None, n_input))
h_hidden = Dense(n_hidden, activation='tanh')(x_input)
y_output = Dense(n_output, activation='linear')(h_hidden)
model = Model(x_input, y_output)
```



# Update parameters using tf.GradientTape() and optimizer

```
def fit(x_trn, y_trn, x_val, y_val, epochs, batch_size):
    trn_loss = []
    val_loss = []
    for epoch in range(epochs):
        # Training with mini-batch
        for batch in range(int(x.shape[0] / batch_size)):
            idx = np.random.choice(x_trn.shape[0], batch_size)
            x_bat = x_trn[idx]
            y_bat = y_trn[idx]

            # Automatic differentiation
            with tf.GradientTape() as tape:
                loss = mse(y_bat, model(x_bat))

            # Find the gradients of loss w.r.t the parameters
            grads = tape.gradient(loss, model.trainable_variables)

            # update parameters by optimizer
            opt.apply_gradients(zip(grads,
                                   model.trainable_variables))

        # loss history
        loss = mse(y_trn, model(x_trn))
        trn_loss.append(loss.numpy())

        loss = mse(y_val, model(x_val))
        val_loss.append(loss.numpy())

    if epoch % 10 == 0:
        print("{:} train_loss={:.4f}, val_loss={:.4f}.\n"
              .format(epoch, trn_loss[-1], val_loss[-1]))

    return trn_loss, val_loss
```



## ■ 3. Tensorflow + Keras functional API – [Example] Nonlinear Regression



```
# training
trn_loss, val_loss = fit(x_train, y_train, x_test, y_test,
                        epochs=200, batch_size=50)

# Visually see the data.
plt.figure(figsize=(7,5))
plt.scatter(x_train, y_train, s=50, c='blue', alpha=0.5,
            label='train')
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.5,
            label='valid')
plt.legend()
plt.show()

# Visually see the loss history.
plt.plot(trn_loss, c='blue', label='train loss')
plt.plot(val_loss, c='red', label='validation loss')
plt.legend()
plt.show()

# Visually check the prediction result.
y_pred = model.predict(x_pred)
plt.figure(figsize=(7,5))
plt.scatter(x_train, y_train, s=20, c='blue', alpha=0.3,
            label='train')
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.3,
            label='validation')
plt.scatter(x_pred, y_pred, s=5, c='red', label='test')
plt.legend()
plt.show()
```

```
0: train_loss=2.2578, val_loss=2.2365
10: train_loss=1.3973, val_loss=1.5123
20: train_loss=1.3752, val_loss=1.4882
30: train_loss=1.3449, val_loss=1.4547
40: train_loss=1.2845, val_loss=1.3782
50: train_loss=1.1699, val_loss=1.2344
60: train_loss=0.9798, val_loss=1.0186
70: train_loss=0.8442, val_loss=0.8505
80: train_loss=0.7464, val_loss=0.7175
90: train_loss=0.7233, val_loss=0.6700
100: train_loss=0.7139, val_loss=0.6551
110: train_loss=0.7140, val_loss=0.6500
120: train_loss=0.7286, val_loss=0.6685
130: train_loss=0.7109, val_loss=0.6508
140: train_loss=0.7482, val_loss=0.6941
150: train_loss=0.7093, val_loss=0.6477
160: train_loss=0.7108, val_loss=0.6534
170: train_loss=0.7356, val_loss=0.6824
180: train_loss=0.7054, val_loss=0.6454
190: train_loss=0.7078, val_loss=0.6485
```





## ■ 4. Custom Loss : Regularized loss function – [Example] Nonlinear Regression

# [MXDL-4-02] 7.custom\_loss.py - Regularized loss function

```
import numpy as np
import tensorflow as tf
from sklearn.datasets import make_blobs
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras import optimizers
import matplotlib.pyplot as plt
```

# Generate a data set

```
x = np.random.random((1000, 1))
y = 2.0 * np.sin(2.0 * np.pi * x) + \
    np.random.normal(0.0, 0.8, (1000, 1))
```

# Generate training, test data set

```
x_train, x_test, y_train, y_test = train_test_split(x, y)
x_pred = np.linspace(0, 1, 200).reshape(-1, 1)
```

# Custom loss: Applying L2 regularization to the loss function

```
class regularized_loss(tf.keras.losses.Loss):
```

```
    def __init__(self, C, h_layer, o_layer):
```

```
        super(regularized_loss, self).__init__()
```

```
        self.C = C
```

```
        self.h_layer=h_layer
```

```
        self.o_layer=o_layer
```

$$L_{reg} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 + C \left( \sum_m w_{h,m}^2 + \sum_n w_{o,n}^2 \right)$$

```
    def call(self, y_true, y_pred):
```

```
        mse = tf.reduce_mean(tf.math.square(y_true - y_pred))
```


```
        wh = self.h_layer.weights[0] # weights in hidden layer
```

```
        wo = self.o_layer.weights[0] # weights in output layer
```

```
        mse += self.C * tf.reduce_sum(tf.math.square(wh))
```

```
        mse += self.C * tf.reduce_sum(tf.math.square(wo))
```

```
        return mse
```



# Create an ANN model

```
n_input = x.shape[1] # number of input neurons
n_output = 1 # number of output neurons
n_hidden = 8 # number of hidden neurons
adam = optimizers.Adam(learning_rate=0.01)
```

```
h_layer = Dense(n_hidden, activation='tanh') # hidden layer
```

```
o_layer = Dense(n_output, activation='linear') # output layer
```

```
x_input = Input(batch_shape=(None, n_input))
```

```
h = h_layer(x_input)
```

```
y_output = o_layer(h)
```

```
model = Model(x_input, y_output)
```

```
myloss = regularized_loss(0.001, h_layer, o_layer)
```

```
model.compile(loss=myloss, optimizer=adam)
```

# Training

```
f = model.fit(x_train, y_train,
              validation_data=(x_test, y_test),
              epochs=200, batch_size=50)
```

# Visually see the data.

```
plt.figure(figsize=(7,5))
```

```
plt.scatter(x_train, y_train, s=50, c='blue', alpha=0.5, label='train')
```

```
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.5, label='valid')
```

```
plt.legend()
```

```
plt.show()
```


# Visually see the loss history.

```
plt.plot(f.history['loss'], c='blue', label='train loss')
```

```
plt.plot(f.history['val_loss'], c='red', label='validation loss')
```

```
plt.legend()
```

```
plt.show()
```



## ■ 4. Custom Loss : Regularized loss function – [Example] Nonlinear Regression

# Visually see the prediction result.

```

y_pred = model.predict(x_pred)
plt.figure(figsize=(7,5))
plt.scatter(x_train, y_train, s=20, c='blue', alpha=0.3,
label='train')
plt.scatter(x_test, y_test, s=20, c='red', alpha=0.3,
label='validation')
plt.scatter(x_pred, y_pred, s=5, c='red', label='test')
plt.legend()
plt.show()

```


[ C = 0.001 ]

```


Epoch 198/200
15/15 [=====] - 0s 3ms/step - loss: 0.7438 - val_loss: 0.7287
Epoch 199/200
15/15 [=====] - 0s 3ms/step - loss: 0.7360 - val_loss: 0.7028
Epoch 200/200
15/15 [=====] - 0s 2ms/step - loss: 0.7381 - val_loss: 0.6930

```

Loss history



Prediction



$$L_{reg} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 + C \left( \sum_m w_{h,m}^2 + \sum_n w_{o,n}^2 \right)$$


[ C = 0.01 ]

```

Epoch 198/200
15/15 [=====] - 0s 2ms/step - loss: 1.1684 - val_loss: 1.1273
Epoch 199/200
15/15 [=====] - 0s 3ms/step - loss: 1.1692 - val_loss: 1.1374
Epoch 200/200
15/15 [=====] - 0s 2ms/step - loss: 1.1658 - val_loss: 1.1238

```

Loss history



Prediction

